



Scan to try TrustLayer live

TRUSTLAYER

Design Thinking Project 2026

1

USER / PROTAGONIST

Our users are **people who rely on AI tools for academic, professional, healthcare, and everyday decision making**. They reach for AI because it is fast, fluent, and convenient.

But the moment accuracy, ethics, safety, or accountability enter the picture, they hesitate. **Students** wonder if AI counts as help or cheating. **Professionals** spot polished outputs hiding subtle errors. **Healthcare workers** avoid AI in critical moments. **Everyday users** fall back on Google to double check.

They need a clearer way to tell **when AI output is trustworthy, when it needs verification, and when it should not be used without human judgement**.

2

CONTEXT & SITUATION

The problem shows up the moment a task starts to **matter**, when correctness, responsibility, or ethical judgement are on the line.

A student under deadline pressure. A professional racing to ship a report. A healthcare worker weighing decision support. An everyday user looking up quick advice.

In every situation, the AI gives a **fast, confident, fluent answer**. The user is left holding the harder question: *is this actually accurate, acceptable, and safe to act on?*

3

EVIDENCE BASE

We grounded the problem in primary research: **user interviews, empathy mapping, 5 Whys analysis, and causal mapping**.

Across every user type, the same pattern surfaced. **People find AI useful, but they do not fully trust it**.

Students felt unsure whether AI use counted as help or cheating. Professionals noticed polished outputs with hidden errors. Healthcare users avoided AI for high stakes calls because mistakes could harm someone. Everyday users habitually cross checked AI answers on Google because they had no other way to verify.

The polish of an AI answer kept rising. The user's confidence in it did not.

4

PAIN POINTS & NEEDS

Users carry the cost of **uncertainty, missing trust, and extra verification work**. AI sounds confident even when it is wrong, incomplete, or unsafe, and offers no signal about which is which.

Emotionally, that lands as **guilt, hesitation, skepticism, and the fear of acting on something wrong**. Users have no easy way to see sources, uncertainty, ethical boundaries, or who is responsible if the answer fails.

What they need: AI tools that make **trust, limits, and verification legible** before they depend on the output, not after.

5

ROOT CAUSES (5 WHYS)

The root cause is structural. Most AI tools are designed to produce **fast, fluent, confident answers**, not to show how trustworthy, safe, or accountable those answers actually are.

The result: users receive polished outputs without visible **evidence, uncertainty, reasoning, or safe-use boundaries**. The burden of judgement quietly shifts from the system to the person.

The user is left to decide, often alone and under time pressure, whether the output is reliable enough to act on. That is the gap our solution fills.

6

HOW MIGHT WE...?

01 How might we surface AI uncertainty **without breaking the user's flow**?

02 How might we give every claim a **visible, checkable source** instead of hiding evidence?

03 How might we score trust at the **sentence level**, not just per response?

04 How might we keep users **actively thinking** rather than passively accepting AI output?

05 How might we make verification **fast enough that users actually do it**?

06 How might we make AI reasoning **inspectable**, not a black box?

07 How might we adapt trust signals to the **stakes and domain** of the question?

7

TOP IDEAS · OUR SOLUTION: TRUSTLAYER

TrustLayer is a paste and analyse tool that scores any AI-generated text **sentence by sentence**. Instead of treating every AI response as equally reliable, it surfaces a confidence signal on each claim so users know exactly how much weight to give it.

Paste output from ChatGPT, Gemini, Claude, or anywhere else. Every sentence is scored instantly. **No API, no backend, no data leaves the browser**.

Each sentence also receives a claim-type tag (sourced reference, sensitive domain, inference, opinion, factual claim, unverified claim) so users see *why* the score is what it is.

High confidence

Well-established fact or cited source. Safe to act on.

Medium confidence

Sensitive domain or hedged language. Verify before acting.

Low confidence

Speculative or future-facing. Do not act without verification.

Trust in AI is not a single feature. It is a system. TrustLayer shifts AI from a black-box answer generator into a transparent, checkable, accountable partner.

PASTE HERE

Ibuprofen 200-400mg every 4-6 hours.
Remote work will dominate by 2030.
Water boils at 100°C at sea level.

ANALYSIS

Ibuprofen 200-400mg... SENSITIVE

Remote work will dominate... INFERENCE

Water boils at 100°C... FACTUAL

Paste

Sentence wise

Real time

Safe

Rule based engine

100% client side

8

PROTOTYPE / FLOW DIAGRAM

A **transparent rule-based heuristic engine**. No ML, no external API. Every sentence runs through four ordered signal checks, and the live decision tree is shown right in the UI. **No black box**.

STEP 1 Cites a named source, study, or URL? **HIGH**

STEP 2 Speculative or future-facing language? **LOW**

STEP 3 Medical, legal, financial, or safety domain? **MEDIUM**

STEP 4 Hedging? Else default to unverified. **MEDIUM**

*We default to **Medium**, not High. A false positive is far less harmful than a false negative, so we err on the side of caution.*

B.Tech
2nd Year

Group Details

Charvi · Mohana · Gurpreet · Lalith · Aneesh · Karthik · Pushpashree · Vaishnavi
B.Tech 2nd Year · Design Thinking Workshop · Final Submission